

4. Partial differential equations

In these notes we discuss how to solve numerically partial differential equations (PDEs), with the method of finite difference. Problems with partial differential equations are solved differently according to whether they are *initial value problems*, where the initial conditions (say at $t = 0$ is specified), and *boundary value problem*, where the values on some boundary are specified. We will consider these two main classes of PDE problem below, focussing on some concrete examples in both cases.

4.1 Initial value problems: discretisation of the diffusion equation

Let us consider the diffusion equation, first in 1 dimension (1D), as our example of initial value problem. The equation we want to solve is therefore

$$\frac{\partial c(x, t)}{\partial t} = D \frac{\partial^2 c(x, t)}{\partial x^2}, \quad (1)$$

subject to the initial condition $c(x, t = 0) = c_0(x)$. Here and in all the discussion of initial value problems, we imagine that there are periodic boundary conditions in space. Finite difference methods are based on discretisations of both space and time. In what follows, we will associate j to the spatial discretisation index, and n to the time discretisation – we will also call $c(j; n)$ the discrete approximation of, say, $c(x, t)$. Two successive points in the spatial discretisation of c are separated by δx , while two time steps are separated by δt . In other words we perform the following mapping

$$\begin{aligned} x &\rightarrow j\delta x \\ t &\rightarrow n\delta t. \end{aligned} \quad (2)$$

We now need to specify how to approximate derivatives by finite differences, i.e. how to discretise the PDE in question.

First, the time derivative can be discretised as follows:

$$\frac{\partial c(x, t)}{\partial t} \simeq \frac{c(x, t + \delta t) - c(x, t)}{\delta t} = \frac{c(j; n + 1) - c(j; n)}{\delta t}, \quad (3)$$

which is known as forward explicit (also Euler) finite difference rule. The reason for the name “explicit” will become more obvious below.

Now we need to discretise spatial derivatives. To this end, let us perform the following Taylor expansions for $\delta x \rightarrow 0$ (similar to what we did when deriving the Verlet algorithm),

$$\begin{aligned} c(x + \delta x, t) &\simeq c(x) + \delta x \frac{\partial c}{\partial x} + \frac{\delta x^2}{2} \frac{\partial^2 c}{\partial x^2} \\ c(x - \delta x, t) &\simeq c(x) - \delta x \frac{\partial c}{\partial x} + \frac{\delta x^2}{2} \frac{\partial^2 c}{\partial x^2}. \end{aligned} \quad (4)$$

By summing the two expansions in Eq. 4, we obtain

$$c(x + \delta x, t) + c(x - \delta x, t) \simeq 2c(x, t) + \delta x^2 \frac{\partial^2 c}{\partial x^2} + \mathcal{O}(\delta x^4) \quad (5)$$

where the error is of order δx^4 because all the odd terms cancel when summing the two expansions in Eq. 4. The discretisation for the (1D) Laplacian is then

$$\frac{\partial^2 c}{\partial x^2} \simeq \frac{c(j+1; n) + c(j-1; n) - 2c(j; n)}{\delta x^2} + \mathcal{O}(\delta x^2), \quad (6)$$

this discretisation is an example of a centred difference rule. Another one is obtained by subtracting out the two expansions in Eq. 4 to obtain an estimate for $\frac{\partial c}{\partial x}$,

$$\frac{\partial c}{\partial x} \simeq \frac{c(j+1; n) - c(j-1; n)}{2\delta x} + \mathcal{O}(\delta x^2). \quad (7)$$

Note that the error is again of order δx^2 .

We can put together the discretisation rules in Eqs. 3 and 6 to obtain the following finite difference approximation for the 1D diffusion equation, Eq. 1,

$$c(j; n+1) = c(j; n) + \frac{D\delta t}{\delta x^2} [c(j+1; n) + c(j-1; n) - 2c(j; n)]. \quad (8)$$

Eq. 8 provides an *explicit* (hence the name of the discretisation in Eq. 3) way to compute the solution at time $n+1$ knowing the solution at time n .

4.2 von Neumann stability analysis

Now we have a finite difference approximation for our PDE and associated initial value problem: it is given by Eq. 8. This is good, but are we sure that the algorithm will be stable, or that it will converge to the exact solution? Intuitively, we expect that if δx and δt are small enough, all should be good. However, it is crucial to be able to answer this question more precisely if we want our PDE algorithm to be of any use in practice!

To understand the stability and convergence properties of our finite difference discretisation, we now perform a calculation, which goes under the name of von Neumann stability analysis. Suppose that our numerical solution is equal to the exact solution, $c_{\text{exact}}(x, t)$, plus a numerical error, $\epsilon(x, t)$: $c(x, t) = c_{\text{exact}}(x, t) + \epsilon(x, t)$. As $c_{\text{exact}}(x, t)$ obeys Eq. 1, its discretised version will obey Eq. 8. Therefore, as $c(j; n)$ also obeys Eq. 8 (it is the definition of our algorithm!), so will the discretisation of the error term, $\epsilon(j; n)$. Therefore, one has

$$\epsilon(j; n+1) = \epsilon(j; n) + \frac{D\delta t}{\delta x^2} [\epsilon(j+1; n) + \epsilon(j-1; n) - 2\epsilon(j; n)]. \quad (9)$$

Let us now expand $\epsilon(x, t)$ in Fourier series

$$\epsilon(x, t) = \sum_m e^{ik_m x} A_m(t) \quad (10)$$

where k_m is an integer multiple of $2\pi/L$, where L is the domain of integration (and equivalently simulation). As the equation is linear, it only suffices to follow the evolution of each of the Fourier component separately. The discretised version of one such component is

$$\epsilon(j; n) = \lambda^n e^{ikj\delta x} \quad (11)$$

where we have assumed $A_m(t) = \lambda^n$: this is as we expect that as $n \rightarrow \infty$ the error either increases (if $\lambda > 1$) or decreases to zero (if $\lambda < 1$). This simple form is appropriate in view of the linearity of the PDE in Eq. 1.

Plugging in the ansatz in Eq. 11 into Eq. 9, we obtain the following “dispersion relation” which determines λ as a function of the wavevector k ,

$$\lambda = 1 + \frac{D\delta t}{\delta x^2} [2 \cos(k\delta x) - 2]. \quad (12)$$

Numerical stability requires the error not to explode with increasing n : this is equivalent to requiring $|\lambda| < 1$. As $0 < [2 - 2 \cos(k\delta x)] < 2$, we obtain that λ is always smaller than 1, so the stability criterion is

$$\lambda = 1 - \frac{D\delta t}{\delta x^2} [2 - 2 \cos(k\delta x)] \geq -1. \quad (13)$$

This needs to hold for any k , otherwise some wavelength will grow exponentially and the numerical error will blow up: therefore one needs $1 - 4 \frac{D\delta t}{\delta x^2} \geq -1$, or equivalently $\frac{D\delta t}{\delta x^2} < \frac{1}{2}$. This classical criterion for the stability of numerical diffusion equations stresses that, if we need to decrease δx to better resolve spatial details in the dynamics, δt needs to be scaled with δx^2 to ensure stability.

4.3 The advection equation

Let us consider as another example the following equation, modelling 1D advection at a constant velocity v ,

$$\frac{\partial c}{\partial t} + v \frac{\partial c}{\partial x} = 0. \quad (14)$$

Starting with an initial condition, $c(x, t = 0) = c_0(x)$, Eq. 14 simply advects it with velocity v , $c(x, t) = c_0(x - vt)$.

By combining Eq. 3 and Eq. 7, we can easily write down the following explicit finite difference discretisation for Eq. 14,

$$c(j; n + 1) = c(j; n) + \frac{v\delta t}{2\delta x} [c(j + 1; n) + c(j - 1; n)]. \quad (15)$$

All seems well: however a stability analysis as in 4.2 reveals a nasty surprise! Using the ansatz in Eq. 11 one obtains the following dispersion relation linking the growth rate of the error, λ , to the wavevector, k ,

$$\lambda = 1 - i \frac{\delta x}{vt} \sin(k\delta x) \quad (16)$$

which leads to a modulus of λ equal to

$$|\lambda| = \sqrt{1 + \left(\frac{v\delta t}{\delta x}\right)^2 \sin^2(k\delta x)} \quad (17)$$

which is always ≥ 1 ! Therefore the simple and intuitive finite difference discretisation in Eq. 15 is always unstable (in jargon we say that the discretisation in Eq. 15 is unconditionally unstable).

In order to avoid this instability, the simplest “fix” is to substitute $c(j; n)$ in the right-hand side of Eq. 15 with $\frac{c(j+1; n) + c(j-1; n)}{2}$, as follows

$$c(j; n+1) = \frac{c(j+1; n) + c(j-1; n)}{2} + \frac{v\delta t}{2\delta x} [c(j+1; n) + c(j-1; n)]. \quad (18)$$

This is known as the Lax prescription. The associated stability criterion can now be found to be $\frac{v\delta t}{\delta x} \leq 1$. The Lax prescription introduces numerical diffusion, therefore if one starts with a sharp kink at $t = 0$, the interface will widen in time. A better choice is to use other discretisation for the advection term $v\frac{\partial c}{\partial x}$, in particular high order (third or more) upwind schemes work very well.

The advection equation provides a simple instructive example which shows how subtle finite difference methods can be! It is often the case that small changes in the discretisation choice, such as Eq. 15 and 18, lead to a very different quality for the resulting algorithm!

4.4 Boundary value problems: the Jacobi relaxation algorithm

We now consider the case of boundary value problems, where the solution is not time-dependent, rather it is specified on the boundary of the simulation domain. A prototype of such problems is the Poisson equation of electrostatics,

$$\nabla^2 \phi = -\frac{\rho}{\epsilon} \quad (19)$$

where ϕ is the electrostatic potential, ρ is the charge density, $\epsilon = \epsilon_r \epsilon_0$ where ϵ_0 and ϵ_r are the permittivity of vacuum and dielectric permittivity respectively. Boundary conditions are specified at the boundary of the simulation domain (in 3D); they can be Dirichlet (specified ϕ), Neumann (specified normal derivative of ϕ), or mixed boundary conditions.

Let us focus for simplicity first on the 1D case. A class of simple, and often used, methods which to solve boundary value problem is that of relaxation methods. These consist in trying to reformulate the boundary value problem as an appropriate initial value problem, which in steady state provides a solution of the problem of interest. In our case, for instance, we may attempt to solve the following initial value problem,

$$\frac{\partial \phi(x, t)}{\partial t} = \frac{\partial^2 \phi(x, t)}{\partial x^2} + \frac{\rho(x)}{\epsilon}. \quad (20)$$

When the dynamics of Eq. 20 reaches steady state, the associated $\phi(x, t)$ provides a valid solution of the boundary value problem in Eq. 19.

A suitable discretisation of Eq. 20 is

$$\phi(j; n+1) = \phi(j; n) + \delta t \frac{[\phi(j+1; n) + \phi(j-1; n) - 2\phi(j; n)]}{\delta x^2} + \delta t \rho(j), \quad (21)$$

where for simplicity we have just taken $\epsilon = 1$. Note that the explicit update for $\phi(j; n+1)$ is computed for j *inside* the simulation domain; whereas every time a point on the boundary is needed (either at timestep n or $n+1$), the appropriate boundary condition is used (for instance, in the case of Dirichlet boundary conditions, we do not evolve the values of ϕ on the boundaries but just set it equal to what they should be there).

Now, we should point out that we are not interested in accurately simulating the dynamics related to Eq. 20: it is anyway fictitious! All we care about is the steady state, equilibrated, solution, which provides our numerical estimate for the boundary value problem which we are really interested in. With this in mind, we choose the largest possible value of δt , namely $\delta t = \frac{\delta x^2}{2}$, to accelerate convergence: we then obtain,

$$\phi(j; n+1) = \frac{1}{2} [\phi(j+1; n) + \phi(j-1; n) + \delta x^2 \rho(j)]. \quad (22)$$

This is the *Jacobi* algorithm. It provides a simple explicit rule to calculate $\phi(j; n+1)$ in terms of value of ϕ at the previous timestep n . This algorithm can also be straightforwardly generalised to more dimensions. For example in 3D the Jacobi algorithm is (a position in 3D is labelled by three indices, i, j and k),

$$\begin{aligned} \phi(i, j, k; n+1) &= \frac{1}{6} \left[\phi(i+1, j, k; n) + \phi(i-1, j, k; n) \right. \\ &+ \phi(i, j+1, k; n) + \phi(i, j-1, k; n) \\ &+ \left. \phi(i, j, k+1; n) + \phi(i, j, k-1; n) + \delta x^2 \rho(i, j, k) \right]. \end{aligned} \quad (23)$$

The Jacobi algorithm is nice and simple; it can be coded up easily and it is therefore quite popular. Its problem is that its convergence is quite slow: one can prove that in order to reduce the numerical error in the steady state estimate for ϕ by a factor of 10^{-p} , one needs an extra pN^2 iterations, where N is the linear dimension of the discretisation lattice (which in D dimensions would thus consist of N^D points). We will now see how to (slightly) modify the Jacobi algorithm to achieve faster convergence to the steady state solution.

4.5 Boundary value problems: the Gauss-Seidel and successive over-relaxation methods

One simple modification of the Jacobi algorithm comes from the realisation that when n is large it would be theoretically equivalent in Eqs. 22, 23 if, instead of the value at timestep n , we used the value of timestep $n+1$ for ϕ on the right

hand side. This leads to the Gauss-Seidel algorithm, which generalises Eq. 22 to

$$\phi(j; n+1) = \frac{1}{2} [\phi(j+1; n) + \phi(j-1; n+1) + \delta x^2 \rho(j)]. \quad (24)$$

where you should note the $\phi(j-1; n+1)$ on the right hand side. This algorithm assumes that we are looping over j from 0 to $N-1$ when updating ϕ . If this is done, then when updating site j at timestep $n+1$ we know $j+1$ only at timestep n , but we know $j-1$ at both timestep n and $n+1$ (as we have just updated $j-1$!) Therefore one may imagine that a faster convergence can be attained if we plug in the *latest* estimate for $\phi(j-1)$ on the right hand side. This is indeed correct (although not so easy to prove!), and the Gauss-Seidel algorithm converges twice as fast as the Jacobi algorithm. The Gauss-Seidel algorithm also allows us to only keep track of one array for ϕ , with only its latest value, as the updated values of $\phi(j)$ are used immediately when updating $\phi(j+1)$.

A problem of the Gauss-Seidel algorithm, which is relevant to high performance computing, is that it is not easy to parallelise, as in common parallel architectures we are not sure when a site is updated with respect to another one. The workaround is to divide the lattice in alternating black and white nodes (like a checkerboard in 2D): then one can update first all black, then all whites etc, without worrying about conflicting/simultaneous updates.

While the Gauss-Seidel algorithm is faster, this is still not enough for applications to real-world problem, for instance to situations where charges can move, so that the electrostatic potential needs to be recomputed very often. A trick to achieve even faster convergence is to use the *successive over-relaxation (sor) method*. This is based on the observation that the typical reason why relaxation methods are slow is that the field/variable under investigation changes less and less with n .

Let us first consider a simple relaxation methods, where the target variable, x , is a scalar rather than a field (our electrostatic potential). A relaxation method then provides a recursion relation to go from $x(n)$, the n -th estimate of x , to $x(n+1)$, say given by $x(n+1) = f(x(n))$ (f is a generic function). Let us now introduce a “relaxation” parameter ω , and define a new recursion as follows

$$x(n+1) = \omega f(x(n)) + (1-\omega)x(n) = f(x(n)) + (\omega-1)\delta x(n) \quad (25)$$

where $\delta x(n) = f(x(n)) - x(n)$. The case $\omega = 0$ corresponds to no relaxation; the case $\omega = 1$ to the $x(n+1) = f(x(n))$ recursion; while $0 < \omega < 1$ and $1 < \omega < 2$ correspond respectively to “under-relax” and “over-relax”, as $x(n+1)$ undershoots and overshoots $f(x(n))$ respectively. Typically there is an optimal ω which renders convergence as fast as possible within this family of models; usually this is in the range $\omega \sim 1.2 - 1.4$ – but it is problem dependent!

In the boundary value problem we are considering, the analog of $f(x(n))$ is the Gauss-Seidel estimate. The SOR model then is equivalent to the following update rule (in 1D)

$$\phi(j; n+1) = (1-\omega)\phi(j; n) + \omega\phi_{GS}(j; n+1) \quad (26)$$

where $\phi_{GS}(j; n + 1)$ denotes the Gauss-Seidel update rule, Eq. 22. The SOR model can be generalised straightforwardly in 3D. You should play around with ω in your code to see what gives the fastest convergence in your case!